




(a)



(b)



(a) covering algorithm; (b) decision tree

- | | | |
|-------------------|------|--------------------------------|
| age = young | 2/8 | (are correct for hard) |
| age = presbyopic | 1/8 | |
| astigmatism = yes | 4/12 | |



continue covering, rule refining...

if astigmatism = yes then recommendation = hard

covers only 4 cases....

if astigmatism = yes and ? then recommendation = hard

age = young 2/4 (are still correct for hard)

age = presbyopic 1/4

tear production rate = normal 4/6

if astigmatism = yes and tear production rate = normal
then recommendation = hard

.....

now consider rules for next class



a simple rule learner (separate-and-conquer)

for each class c

`e := {full set of instances};`

`while e contains instances in c`

`create rule r with empty LHS, predicting c;`

`while r is not perfect and there are attrs`

`for each attr a not in r, and each val v,`

`consider adding a=v to LHS of r;`

`select a, v to maximize accuracy pos/tot;`

`(break ties by picking cond with max pos)`

`add a=v to r;`

`remove instances covered by r from e;`

rules for each class need not be considered in order although instances covered by a rule are removed when rule is completed...



criteria for choosing a test to be added

- *pos/tot*
 - maximizes rule correctness
 - max: no neg examples covered
 - prefer test covering 1 pos, 0 neg, over test with 100 pos, 1 neg
 - finds & eliminates special cases first
- info gain measure (accuracy increase)
 - $pos [\log(pos/tot) - \log(Pos/Tot)]$
 - Pos, Tot*: # of data covered by rule **before** new test added
 - maximizes pos coverage, minimizes neg coverage
 - finds hi-coverage, general rules first, exceptions later



making sensible rules - avoid overfitting

if astigmatism = yes and tear prod rate = normal then *hard*

correct 4 / 6 cases; default rule (always recommends *hard*): 4 / 24;
so rule greatly improves accuracy, with small (0.14%) prob of
improvement being due to chance

if astigmatism = yes then *hard*

improves accuracy from 4/24 to 4/12, with 4.7% prob of
improvement being due to chance (not a very good rule)

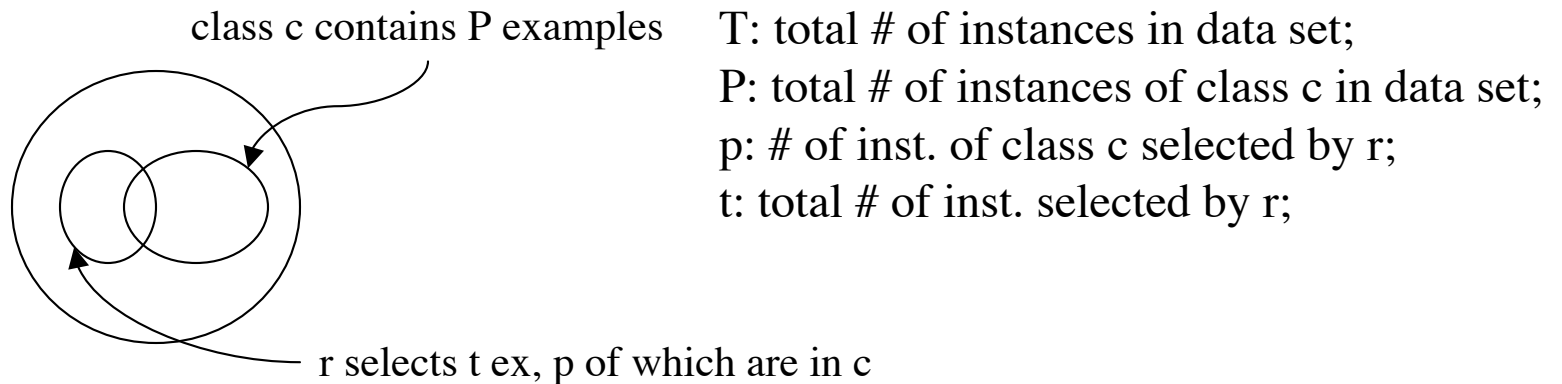
1. make a perfect rule, adding tests until full correctness
2. tentatively remove last test added, see if improvement prob
being due to chance is small; if not, remove test; continue...



decision lists (rule order matters)

```
e := {instances};
while e not empty do
  for each class c for which e has an instance
    use basic covering alg to create best rule for c;
    *compute prob measure  $m(r)$  for rule r, and for rule with last
      test omitted,  $m(r-)$ ;
    while  $m(r-) < m(r)$  do remove last test from rule, repeat *;
    from (best for class) rules generated, pick one with  $\min m(r)$ ;
    output the rule r;
    remove instances covered by rule r from e;
  endwhile
// accuracy measure for growing, prob measure for pruning rules
// does not consider classes one by one but makes rule for every class and picks best one...
//  $m(r-) < m(r)$ : since r- is better (less likely due to chance) than r (with test), just drop the test.
```

probability measure for rule evaluation




e.g. r has accuracy $4/6$ (p/t), default rule has accuracy $4/24$ (P/T).

what's the probability that a random rule **with same coverage** as r has an accuracy improvement as good as (or better than) r ?

use hypergeometric prob distribution
(assume selection without replacement)

prob measure for rule evaluation, cont.

prob [of t cases selected at random, exactly i are in class c] =



$m(r)$ can be used to
compare 2 rules or
find best pruned version

prob that random rule will do as well or better than r is

$$m(r) = \sum_{i=p}^{\min(t,P)} \text{prob [of } t \text{ cases selected at random, exactly } i \text{ are in class } c]$$

small values of $m(r)$ are good: it's unlikely this rule would have happened by chance; there are several cheaper approximations to $m(r)$...

prob measure for rule evaluation, cont.

assuming large T , and sampling with replacement (i.e. assuming constant prob P/T that an instance is in c):

prob [of t cases selected at random, exactly i are in class c] =

$$\binom{t}{i} \left(\frac{P}{T} \right)^i \left(1 - \frac{P}{T} \right)^{t-i}$$

use this in computing

$$m(r) = \sum_{i=p}^{\min(t,P)} \text{prob [of } t \text{ cases selected at random, exactly } i \text{ are in class } c]$$

there are some still cheaper approximations to this approximation to $m(r)$...



incremental reduced-error pruning

split training data into Grow set and Prune set, make a good rule using rule covering, delete a test and try truncated rule on Prune set, see if it's better.... repeat for each class ...

disadvantage: fewer instances available for training, some instances might only be in Prune set, small Prune set might wrongly prefer a rule...

mitigate by **resplitting** into Grow and Prune ...

test rule quality by $m(r)$ or with simpler measures such as:

1. use accuracy p/t . Bad because it prefers rule that gets 1 instance right out of coverage of 1 to rule that gets 1000 right out of 1001
2. use $(p - n)/t$. Same problem.
3. use $(p + (N - n)) / T$ where $n = t - p$ (neg covered), $N = T - P$ (total # of neg instances), $N - n$ not covered neg instances. I.e. how well would rule discriminate predicted class if it were the only rule?
Bad because it treats **noncoverage of neg instances as equally important as coverage of pos instances**; misleading in context of many other rules!

so stick with $m(r)$



incremental reduced-error pruning, cont.

```
e := {instances};
while e is not empty do
{ split e into Grow and Prune in ratio 2:1;
for each class c for which Grow and Prune both contain an instance
  use basic covering algorithm to make best 'perfect' rule for c;
  • compute worth  $w(r)$  for  $r$  and  $w(r-)$  [ $r$  without last test] on Prune;
  while  $w(r-) > w(r)$  remove last condition from  $r$  and repeat •;
from all rules made, select the rule with max  $w(r)$  (or min  $m(r)$ );
output  $r$ ;
e := e - {instances covered by  $r$ };
}
```

// worth could be $1 - m(r)$ or



another method for getting rules

- **combine** DTI and rule covering
- DTI part
 - make a rule by building a **partial** pruned decision tree on current instances
 - it's partial because some branches may be left unexplored
 - turn leaf with largest coverage into **one** rule
 - this may produce a simple and general rule
 - discard tree
 - if data are relatively noise-free, only one path has to be built
- rule covering part
 - as in basic rule covering algorithm, once a rule is made, remove instances covered by it
... until no more instances remain



combining DTI and rule covering

expandSubset(s):

choose test T to split examples into subsets (using info gain);


sort subsets by increasing average entropy (low-entropy subsets may lead to small subtrees and general rules; later subsets may not be expanded ...);

while there is unexpanded subset t' & all expanded subsets are leaves


expandSubset(t'):

if all expanded subsets are leaves & estimated error for subtree \geq error for node, undo expansion into subsets & make node into leaf (i.e. subtree replacement);


after replacement, continue backtracking with siblings of replaced node; if an inner node has all children non-leaves (i.e. when subtree replacement is not done), leave remaining subsets unexplored. From this partial tree, extract a single rule for the leaf that covers most instances




expand 3
because of low
entropy




leaf has lower entropy but can't be
expanded; so backtrack and expand 5:



all children of 5 are
leaves, thus try and
accept pruning:



now consider 3 for
replacement, and accept;
backtrack to expand 4:



suppose 4 is not replaced.
Each leaf could be a rule. We
pick leaf with max coverage
to make **one** rule, then
discard the tree!



labor data: pruned tree vs decision list

```
wage-increase-first-year <= 2.5: bad
wage-increase-first-year > 2.5
|   statutory-holidays <= 10: bad
|   statutory-holidays > 10: good
```

1. wage-increase-first-year > 2.5 AND
longterm-disability-assistance = yes AND
statutory-holidays > 10: good
2. wage-increase-first-year <= 4 AND
working-hours > 36: bad
3. : good



iris data: tree vs decision list

```
petalwidth <= 0.6: Iris-setosa
petalwidth > 0.6
|   petalwidth <= 1.7
|   |   petallength <= 4.9: Iris-versicolor
|   |   petallength > 4.9
|   |   |   petalwidth <= 1.5: Iris-virginica
|   |   |   petalwidth > 1.5: Iris-versicolor
|   petalwidth > 1.7: Iris-virginica
```

1. petalwidth <= 0.6: Iris-setosa
2. petalwidth <= 1.7 AND
petallength <= 4.9: Iris-versicolor
3. : Iris-virginica